Step 1 – Go to address https://54.87.253.194:8888, and ignore warnings.
The IP is permanent, so you can bookmark this site.

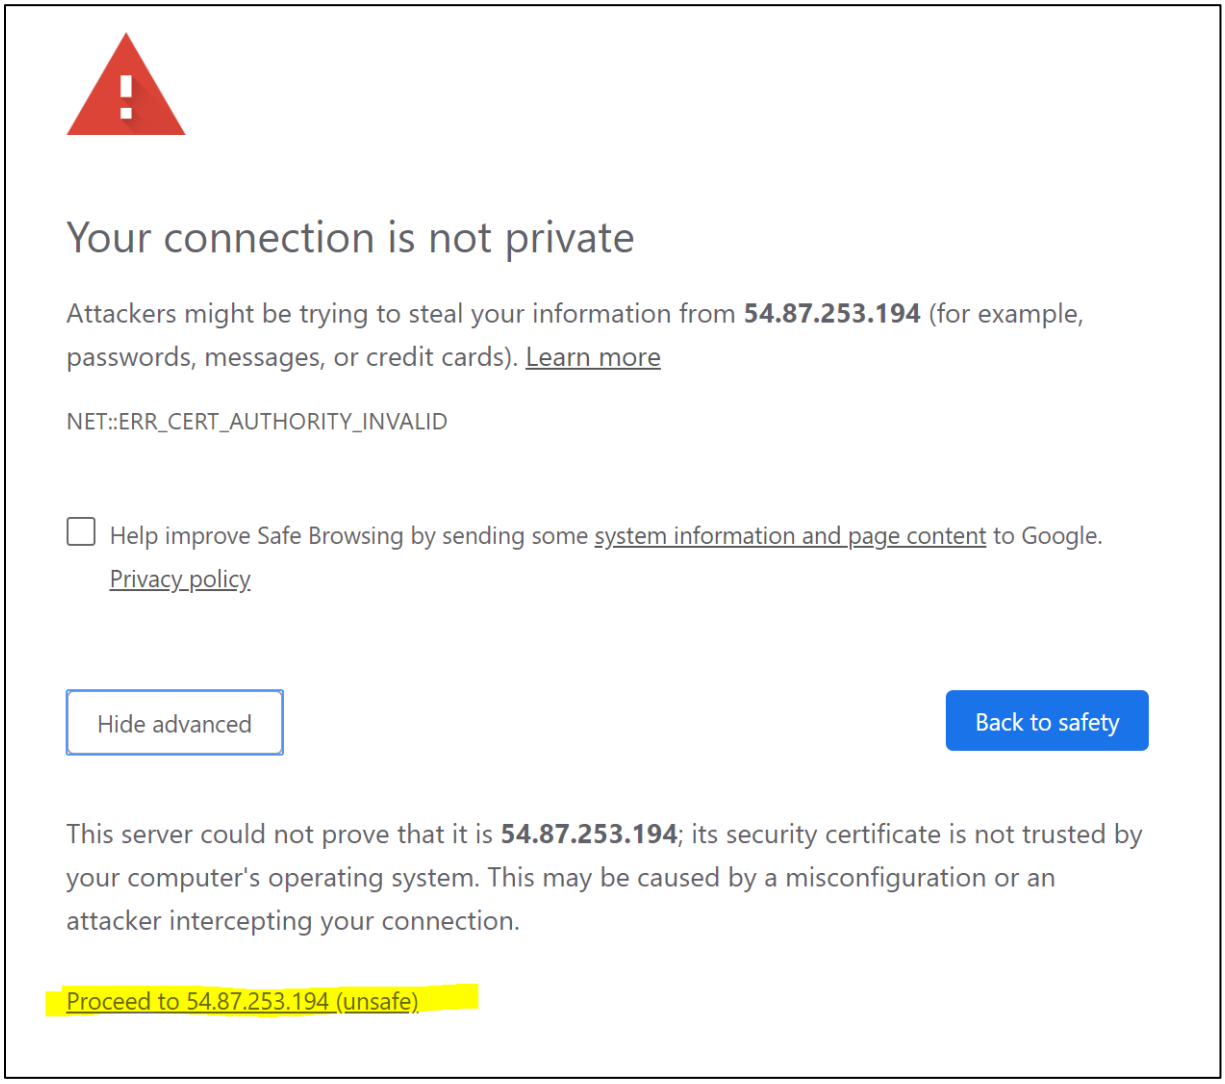

Step 2 – enter password: AI2019

Step 3 – don't change anything, go to folder "AI Marketing", then create a new folder with your name. You can store your file in your folder

Note:

If any package is not available, I will install it from backend. If you need to unzip file, I will do it from backend as well.

We are currently using instance “p2.xlarge”, but we can easily upscale it latter. Together, we are sharing 75G storage, 61G memory, 1 GPU with ~11G memory, and 4 CPU.

With this instance, I have no problem to have train_generator (batch_size=231) and validation_generator (batch_size=143). I believe this is our allowed max setting.

P2 Instance Details

Name	GPUs	vCPUs	RAM (GiB)	Network Bandwidth	Price/Hour*	RI Price / Hour**
p2.xlarge	1	4	61	High	$0.900	$0.425
p2.8xlarge	8	32	488	10 Gbps	$7.200	$3.400
p2.16xlarge	16	64	732	20 Gbps	$14.400	$6.800

P2 instances provide up to 16 **NVIDIA K80 GPUs**, 64 vCPUs and 732 GiB of host memory, with a combined 192 GB of GPU memory, 40 thousand parallel processing cores, 70 teraflops of single precision floating point performance, and over 23 teraflops of double precision floating point performance. P2 instances also offer GPUDirect™ (peer-to-peer GPU communication) capabilities for up to 16 GPUs, so that multiple GPUs can work together within a single host.